

Final Assignment

February 9, 2023

Extracting and Visualizing Stock Data

Description

Extracting essential data from a dataset and displaying it is a necessary part of data science; therefore individuals can make correct decisions based on the data. In this assignment, you will extract some stock data, you will then display this data in a graph.

Table of Contents

- Define a Function that Makes a Graph
- Question 1: Use yfinance to Extract Stock Data
- Question 2: Use Webscraping to Extract Tesla Revenue Data
- Question 3: Use yfinance to Extract Stock Data
- Question 4: Use Webscraping to Extract GME Revenue Data
- Question 5: Plot Tesla Stock Graph
- Question 6: Plot GameStop Stock Graph

Estimated Time Needed: 30 min

```
[1]: !pip install yfinance==0.1.67
!mamba install bs4==4.10.0 -y
!pip install nbformat==4.2.0
```

Collecting yfinance==0.1.67

Downloading yfinance-0.1.67-py2.py3-none-any.whl (25 kB)

Requirement already satisfied: pandas>=0.24 in

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from yfinance==0.1.67) (1.3.5)

Requirement already satisfied: requests>=2.20 in

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from yfinance==0.1.67) (2.28.1)

Requirement already satisfied: lxml>=4.5.1 in

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from yfinance==0.1.67) (4.6.4)

Requirement already satisfied: multitasking>=0.0.7 in

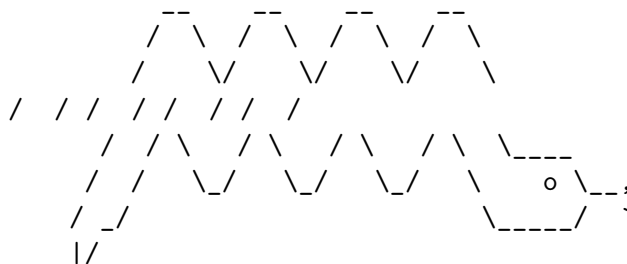
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from yfinance==0.1.67) (0.0.11)

Requirement already satisfied: numpy>=1.15 in

```

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
yfinance==0.1.67) (1.21.6)
Requirement already satisfied: python-dateutil>=2.7.3 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
pandas>=0.24->yfinance==0.1.67) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
pandas>=0.24->yfinance==0.1.67) (2022.6)
Requirement already satisfied: charset-normalizer<3,>=2 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
requests>=2.20->yfinance==0.1.67) (2.1.1)
Requirement already satisfied: certifi>=2017.4.17 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
requests>=2.20->yfinance==0.1.67) (2022.12.7)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
requests>=2.20->yfinance==0.1.67) (1.26.13)
Requirement already satisfied: idna<4,>=2.5 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
requests>=2.20->yfinance==0.1.67) (3.4)
Requirement already satisfied: six>=1.5 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from python-
dateutil>=2.7.3->pandas>=0.24->yfinance==0.1.67) (1.16.0)
Installing collected packages: yfinance
  Attempting uninstall: yfinance
    Found existing installation: yfinance 0.2.4
    Uninstalling yfinance-0.2.4:
      Successfully uninstalled yfinance-0.2.4
Successfully installed yfinance-0.1.67

```



mamba (0.15.3) supported by @QuantStack

GitHub: <https://github.com/mamba-org/mamba>

Twitter: <https://twitter.com/QuantStack>

Looking for: ['bs4==4.10.0']

```
pkgs/r/linux-64      [>                ] (--:-- ) No change
pkgs/r/linux-64      [=====] (00m:00s) No change
pkgs/main/linux-64   [>                ] (--:-- ) No change
pkgs/main/linux-64   [=====] (00m:00s) No change
pkgs/r/noarch        [>                ] (--:-- ) No change
pkgs/r/noarch        [=====] (00m:00s) No change
pkgs/main/noarch     [>                ] (--:-- ) No change
pkgs/main/noarch     [=====] (00m:00s) No change
```

Pinned packages:

- python 3.7.*

Transaction

Prefix: /home/jupyterlab/conda/envs/python

All requested packages already installed

Collecting nbformat==4.2.0

Downloading nbformat-4.2.0-py2.py3-none-any.whl (153 kB)

153.3/153.3 kB

15.3 MB/s eta 0:00:00

Requirement already satisfied: jupyter-core in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
nbformat==4.2.0) (4.12.0)

Requirement already satisfied: traitlets>=4.1 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
nbformat==4.2.0) (5.6.0)

Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
nbformat==4.2.0) (4.17.3)

Requirement already satisfied: ipython-genutils in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
nbformat==4.2.0) (0.2.0)

Requirement already satisfied: importlib-resources>=1.4.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (5.10.1)

Requirement already satisfied: attrs>=17.4.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (22.1.0)

```

Requirement already satisfied: pkgutil-resolve-name>=1.3.10 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (1.3.10)
Requirement already satisfied: typing-extensions in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (4.4.0)
Requirement already satisfied: importlib-metadata in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (4.11.4)
Requirement already satisfied: pyrsistent!=0.17.0,!0.17.1,!0.17.2,>=0.14.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (0.19.2)
Requirement already satisfied: zipp>=3.1.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from importlib-
resources>=1.4.0->jsonschema!=2.5.0,>=2.4->nbformat==4.2.0) (3.11.0)
Installing collected packages: nbformat
  Attempting uninstall: nbformat
    Found existing installation: nbformat 5.7.0
    Uninstalling nbformat-5.7.0:
      Successfully uninstalled nbformat-5.7.0
ERROR: pip's dependency resolver does not currently take into account all
the packages that are installed. This behaviour is the source of the following
dependency conflicts.

nbconvert 7.2.6 requires nbformat>=5.1, but you have nbformat 4.2.0 which is
incompatible.

nbclient 0.7.2 requires nbformat>=5.1, but you have nbformat 4.2.0 which is
incompatible.

jupyter-server 1.23.3 requires nbformat>=5.2.0, but you have nbformat 4.2.0
which is incompatible.

Successfully installed nbformat-4.2.0

```

```

[2]: import yfinance as yf
import pandas as pd
import requests
from bs4 import BeautifulSoup
import plotly.graph_objects as go
from plotly.subplots import make_subplots

```

0.1 Define Graphing Function

In this section, we define the function `make_graph`. You don't have to know how the function works, you should only care about the inputs. It takes a dataframe with stock data (dataframe must contain Date and Close columns), a dataframe with revenue data (dataframe must contain

Date and Revenue columns), and the name of the stock.

```
[3]: def make_graph(stock_data, revenue_data, stock):
    fig = make_subplots(rows=2, cols=1, shared_xaxes=True,
    ↪ subplot_titles=("Historical Share Price", "Historical Revenue"),
    ↪ vertical_spacing = .3)
    stock_data_specific = stock_data[stock_data.Date <= '2021--06-14']
    revenue_data_specific = revenue_data[revenue_data.Date <= '2021-04-30']
    fig.add_trace(go.Scatter(x=pd.to_datetime(stock_data_specific.Date,
    ↪ infer_datetime_format=True), y=stock_data_specific.Close.astype("float"),
    ↪ name="Share Price"), row=1, col=1)
    fig.add_trace(go.Scatter(x=pd.to_datetime(revenue_data_specific.Date,
    ↪ infer_datetime_format=True), y=revenue_data_specific.Revenue.
    ↪ astype("float"), name="Revenue"), row=2, col=1)
    fig.update_xaxes(title_text="Date", row=1, col=1)
    fig.update_xaxes(title_text="Date", row=2, col=1)
    fig.update_yaxes(title_text="Price ($US)", row=1, col=1)
    fig.update_yaxes(title_text="Revenue ($US Millions)", row=2, col=1)
    fig.update_layout(showlegend=False,
    height=900,
    title=stock,
    xaxis_rangeslider_visible=True)
    fig.show()
```

0.2 Question 1: Use yfinance to Extract Stock Data

Using the Ticker function enter the ticker symbol of the stock we want to extract data on to create a ticker object. The stock is Tesla and its ticker symbol is TSLA.

```
[10]: ticker = yf.Ticker("TSLA")
```

Using the ticker object and the function history extract stock information and save it in a dataframe named tesla_data. Set the period parameter to max so we get information for the maximum amount of time.

```
[22]: tesla_data = ticker.history(period="max")
tesla_data
```

```
[22]:
```

	Open	High	Low	Close	Volume \
Date					
2010-06-29	1.266667	1.666667	1.169333	1.592667	281494500
2010-06-30	1.719333	2.028000	1.553333	1.588667	257806500
2010-07-01	1.666667	1.728000	1.351333	1.464000	123282000
2010-07-02	1.533333	1.540000	1.247333	1.280000	77097000
2010-07-06	1.333333	1.333333	1.055333	1.074000	103003500
...
2023-02-02	187.330002	196.750000	182.610001	188.270004	217448300
2023-02-03	183.949997	199.000000	183.690002	189.979996	231684200

2023-02-06	193.009995	198.169998	189.919998	194.759995	186188100
2023-02-07	196.429993	197.500000	189.550003	196.809998	186010300
2023-02-08	196.100006	203.000000	194.309998	201.289993	180154500

	Dividends	Stock Splits
Date		
2010-06-29	0	0.0
2010-06-30	0	0.0
2010-07-01	0	0.0
2010-07-02	0	0.0
2010-07-06	0	0.0
...
2023-02-02	0	0.0
2023-02-03	0	0.0
2023-02-06	0	0.0
2023-02-07	0	0.0
2023-02-08	0	0.0

[3176 rows x 7 columns]

Reset the index using the `reset_index(inplace=True)` function on the `tesla_data` DataFrame and display the first five rows of the `tesla_data` dataframe using the `head` function. Take a screenshot of the results and code from the beginning of Question 1 to the results below.

```
[24]: #tesla_data.set_index('Date', inplace=True)
tesla_data.reset_index(inplace=True)
tesla_data.head()
```

```
[24]:   index      Date      Open      High      Low      Close      Volume \
0      0 2010-06-29  1.266667  1.666667  1.169333  1.592667  281494500
1      1 2010-06-30  1.719333  2.028000  1.553333  1.588667  257806500
2      2 2010-07-01  1.666667  1.728000  1.351333  1.464000  123282000
3      3 2010-07-02  1.533333  1.540000  1.247333  1.280000   77097000
4      4 2010-07-06  1.333333  1.333333  1.055333  1.074000  103003500
```

	Dividends	Stock Splits
0	0	0.0
1	0	0.0
2	0	0.0
3	0	0.0
4	0	0.0

0.3 Question 2: Use Webscraping to Extract Tesla Revenue Data

Use the `requests` library to download the webpage <https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-PY0220EN-SkillsNetwork/labs/project/revenue.htm> Save the text of the response as a variable named `html_data`.

```
[27]: url = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
↳IBMDeveloperSkillsNetwork-PY0220EN-SkillsNetwork/labs/project/revenue.htm"
response = requests.get(url)
html_data = response.text
```

Parse the html data using beautiful_soup.

```
[28]: soup = BeautifulSoup(html_data, 'html.parser')
```

Using BeautifulSoup or the read_html function extract the table with Tesla Quarterly Revenue and store it into a dataframe named tesla_revenue. The dataframe should have columns Date and Revenue.

[Click here](#) if you need help locating the table

Below is the code to isolate the table, you will now need to loop through the rows and columns

```
soup.find_all("tbody")[1]
```

If you want to use the read_html function the table is located at index 1

```
[31]: table = soup.find_all('table')[1]
tesla_revenue = pd.read_html(str(table))[0]
tesla_revenue.columns = ['Date', 'Revenue']
tesla_revenue
```

```
[31]:
```

	Date	Revenue
0	2022-09-30	\$21,454
1	2022-06-30	\$16,934
2	2022-03-31	\$18,756
3	2021-12-31	\$17,719
4	2021-09-30	\$13,757
5	2021-06-30	\$11,958
6	2021-03-31	\$10,389
7	2020-12-31	\$10,744
8	2020-09-30	\$8,771
9	2020-06-30	\$6,036
10	2020-03-31	\$5,985
11	2019-12-31	\$7,384
12	2019-09-30	\$6,303
13	2019-06-30	\$6,350
14	2019-03-31	\$4,541
15	2018-12-31	\$7,226
16	2018-09-30	\$6,824
17	2018-06-30	\$4,002
18	2018-03-31	\$3,409
19	2017-12-31	\$3,288

20	2017-09-30	\$2,985
21	2017-06-30	\$2,790
22	2017-03-31	\$2,696
23	2016-12-31	\$2,285
24	2016-09-30	\$2,298
25	2016-06-30	\$1,270
26	2016-03-31	\$1,147
27	2015-12-31	\$1,214
28	2015-09-30	\$937
29	2015-06-30	\$955
30	2015-03-31	\$940
31	2014-12-31	\$957
32	2014-09-30	\$852
33	2014-06-30	\$769
34	2014-03-31	\$621
35	2013-12-31	\$615
36	2013-09-30	\$431
37	2013-06-30	\$405
38	2013-03-31	\$562
39	2012-12-31	\$306
40	2012-09-30	\$50
41	2012-06-30	\$27
42	2012-03-31	\$30
43	2011-12-31	\$39
44	2011-09-30	\$58
45	2011-06-30	\$58
46	2011-03-31	\$49
47	2010-12-31	\$36
48	2010-09-30	\$31
49	2010-06-30	\$28
50	2010-03-31	\$21
51	2009-12-31	NaN
52	2009-09-30	\$46
53	2009-06-30	\$27

Execute the following line to remove the comma and dollar sign from the **Revenue** column.

```
[32]: tesla_revenue["Revenue"] = tesla_revenue['Revenue'].str.replace(',|\$',"")
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-
packages/ipykernel_launcher.py:1: FutureWarning: The default value of regex will
change from True to False in a future version.
```

```
    """Entry point for launching an IPython kernel.
```

Execute the following lines to remove an null or empty strings in the **Revenue** column.

```
[33]: tesla_revenue.dropna(inplace=True)
```



```
tesla_revenue = tesla_revenue[tesla_revenue['Revenue'] != ""]
```

Display the last 5 row of the `tesla_revenue` dataframe using the `tail` function. Take a screenshot of the results.

```
[35]: tesla_revenue.tail(5)
```

```
[35]:
```

	Date	Revenue
48	2010-09-30	31
49	2010-06-30	28
50	2010-03-31	21
52	2009-09-30	46
53	2009-06-30	27

0.4 Question 3: Use `yfinance` to Extract Stock Data

Using the `Ticker` function enter the ticker symbol of the stock we want to extract data on to create a ticker object. The stock is GameStop and its ticker symbol is `GME`.

```
[38]: gme = yf.Ticker("GME")
gme
```

```
[38]: yfinance.Ticker object <GME>
```

Using the ticker object and the function `history` extract stock information and save it in a dataframe named `gme_data`. Set the `period` parameter to `max` so we get information for the maximum amount of time.

```
[42]: gme_data = gme.history(period="max")
gme_data
```

```
[42]:
```

	Open	High	Low	Close	Volume	Dividends	\
Date							
2002-02-13	1.620128	1.693350	1.603296	1.691666	76216000	0.0	
2002-02-14	1.712707	1.716074	1.670626	1.683250	11021600	0.0	
2002-02-15	1.683250	1.687458	1.658001	1.674834	8389600	0.0	
2002-02-19	1.666418	1.666418	1.578047	1.607504	7410400	0.0	
2002-02-20	1.615921	1.662210	1.603296	1.662210	6892800	0.0	
...		
2023-02-02	22.440001	24.400000	22.219999	22.700001	7655700	0.0	
2023-02-03	22.010000	23.930000	21.799999	22.250000	4930200	0.0	
2023-02-06	21.879999	24.879999	21.770000	23.860001	8708400	0.0	
2023-02-07	23.000000	23.049999	20.500000	21.219999	9093700	0.0	
2023-02-08	21.430000	21.549999	20.610001	21.100000	2841800	0.0	

	Stock Splits
Date	
2002-02-13	0.0

```

2002-02-14      0.0
2002-02-15      0.0
2002-02-19      0.0
2002-02-20      0.0
...
2023-02-02      0.0
2023-02-03      0.0
2023-02-06      0.0
2023-02-07      0.0
2023-02-08      0.0

```

[5284 rows x 7 columns]

Reset the index using the `reset_index(inplace=True)` function on the `gme_data` DataFrame and display the first five rows of the `gme_data` dataframe using the `head` function. Take a screenshot of the results and code from the beginning of Question 3 to the results below.

```
[47]: gme_data.reset_index(inplace=True)
      gme_data.head(5)
```

```
[47]:
```

	level_0	index	Date	Open	High	Low	Close	\
0	0	0	2002-02-13	1.620128	1.693350	1.603296	1.691666	
1	1	1	2002-02-14	1.712707	1.716074	1.670626	1.683250	
2	2	2	2002-02-15	1.683250	1.687458	1.658001	1.674834	
3	3	3	2002-02-19	1.666418	1.666418	1.578047	1.607504	
4	4	4	2002-02-20	1.615921	1.662210	1.603296	1.662210	

	Volume	Dividends	Stock Splits
0	76216000	0.0	0.0
1	11021600	0.0	0.0
2	8389600	0.0	0.0
3	7410400	0.0	0.0
4	6892800	0.0	0.0

0.5 Question 4: Use Webscraping to Extract GME Revenue Data

Use the `requests` library to download the webpage <https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-PY0220EN-SkillsNetwork/labs/project/stock.html>. Save the text of the response as a variable named `html_data`.

```
[53]: url = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
      ↪IBMDeveloperSkillsNetwork-PY0220EN-SkillsNetwork/labs/project/stock.html"
      response = requests.get(url)
      html_data = response.text
      #html_data
```

Parse the html data using `beautiful_soup`.

```
[57]: soup = BeautifulSoup(html_data, 'html.parser')
      #soup
```

Using BeautifulSoup or the `read_html` function extract the table with GameStop Quarterly Revenue and store it into a dataframe named `gme_revenue`. The dataframe should have columns `Date` and `Revenue`. Make sure the comma and dollar sign is removed from the `Revenue` column using a method similar to what you did in Question 2.

[Click here](#) if you need help locating the table

Below is the code to isolate the table, you will now need to loop through the rows and columns

```
soup.find_all("tbody")[1]
```

If you want to use the `read_html` function the table is located at index 1

```
[64]: table = soup.find_all("table")[1]
      gme_revenue = pd.read_html(str(table))[0]
      gme_revenue.columns = ['Date', 'Revenue']

      gme_revenue["Revenue"] = gme_revenue['Revenue'].str.replace(',|\$', "")
      gme_revenue
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-
packages/ipykernel_launcher.py:6: FutureWarning: The default value of regex will
change from True to False in a future version.
```

```
[64]:
```

	Date	Revenue
0	2020-04-30	1021
1	2020-01-31	2194
2	2019-10-31	1439
3	2019-07-31	1286
4	2019-04-30	1548
..
57	2006-01-31	1667
58	2005-10-31	534
59	2005-07-31	416
60	2005-04-30	475
61	2005-01-31	709

```
[62 rows x 2 columns]
```

Display the last five rows of the `gme_revenue` dataframe using the `tail` function. Take a screenshot of the results.

```
[66]: gme_revenue.tail()
```

```
[66]:      Date Revenue
57  2006-01-31    1667
58  2005-10-31     534
59  2005-07-31     416
60  2005-04-30     475
61  2005-01-31     709
```


0.6 Question 5: Plot Tesla Stock Graph

Use the `make_graph` function to graph the Tesla Stock Data, also provide a title for the graph. The structure to call the `make_graph` function is `make_graph(tesla_data, tesla_revenue, 'Tesla')`. Note the graph will only show data upto June 2021.

```
[69]: import matplotlib.pyplot as plt

def make_graph(data,title):
    data.plot(x='Date', y='Close')    #, kind = 'Line')
    plt.title(title)
    plt.xlabel('Date')
    plt.ylabel('Closing Price')
    plt.show()

make_graph(tesla_data, 'Tesla Stock Data')
```




0.7 Question 6: Plot GameStop Stock Graph

Use the `make_graph` function to graph the GameStop Stock Data, also provide a title for the graph. The structure to call the `make_graph` function is `make_graph(gme_data, gme_revenue, 'GameStop')`. Note the graph will only show data upto June 2021.

```
[71]: #def make_graph(data, revenue, title):
#     fig, ax1 = plt.subplots()

def make_graph(data, revenue, title):
    data.plot(x='Date', y='Close')    #, kind = 'Line')
    plt.title(title)
    plt.xlabel('Date')
    plt.ylabel('Closing Price')
    plt.show()

make_graph(gme_data, gme_revenue, 'GME Stock Data')
```



About the Authors:

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Azim Hirjani

0.8 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2022-02-28	1.2	Lakshmi Holla	Changed the URL of GameStop
2020-11-10	1.1	Malika Singla	Deleted the Optional part
2020-08-27	1.0	Malika Singla	Added lab to GitLab

##

© IBM Corporation 2020. All rights reserved.